

Style and Efficiency

Style and Efficiency


- Style
 - Applications and Modules
 - Libraries
 - Return Values
 - Internal Data Structures
 - Concurrency
 - Conventions
- Efficiency

Applications and Modules

- Prefix module names part of a given application or library with the name of the application/library
- If many submodules exist for the same application, create an interface module acting as a single point of entry
 - Allows more flexibility to refactor internal code.
- For the user, the complexity of a module is proportional to the number of exported functions
 - Try to write modules with fewer responsibilities and exported functions


Applications and Modules

Reorganise code!



- Try to reduce inter-module dependencies
 - Makes it easier to refactor
- The call graph should be an acyclic graph

Applications and Modules



- Try to reduce inter-module dependencies
 - Makes it easier to refactor
- The call graph should be an acyclic graph

Libraries

- Collections of related functions should be regrouped into libraries
- The functions should be free of side effects
- If they have side effects, make sure they are related together
 - Functions manipulating a given ETS table, for example
- Document the exported functions!

Return Values

- The return value of a failing function call should be distinct from a successful one
 - it is sometimes impossible to know if there is an error or a value that is the same as the error case
 - tag tuples: **{ok, undefined}** vs. **{error, undefined}**.
 - If the function call always returns a value of a different type when it successful than when it fails, tagging tuples is not necessary
- Pick values that will simplify the caller's task
 - Do not tag values if the call should always succeed, let it crash when it fails



Internal Data Structures

```
-module(q).
-export([add/2, fetch/1]).

add(Item, Q) ->
    lists:append(Q, [Item]).

fetch([]) -> {error, empty}.
fetch([H|T]) -> {ok, H, T};

%% Used as follows
main() ->
    NewQ = [],
    Queue1 = q:add(joe, NewQ),
    Queue2 = q:add(mike, Queue1).
```

- Do not allow private data to leak out
 - All details about the data structure should be abstracted away by the interface
- In the example, **NewQ = []** exposes information about the structure of the queue
 - It is impossible to change the representation without changing all the callers



Internal Data Structures

```
-module(q).
-export([add/2, fetch/1,
        empty/0]).

empty() -> [].

...

%% Used as follows
main() ->
    NewQ = q:empty(),
    Queue1 = q:add(joe, NewQ),
    Queue2 = q:add(mike, Queue1).
```

- The new representation is equivalent in functionality
- Allows to change the queue representation without modifying the callers



Concurrency: hiding information

```
resource_server ! {free, Resource}
%% becomes
resource_server:free(Resource)
```

- Hide all message passing in a **functional interface** for greater flexibility, debugging capabilities and information hiding
- Place the client functions in the same module as the process
 - Makes it easier to follow the message flow without jumping between modules
- Create a mapping between each process and a concurrent activity in the system you are modelling.



Concurrency: tagging messages

```
loop(State) ->
    receive
        {Mod, Fun, Args} ->
            NewState = apply(Mod, Fun,
                             Args),
            loop(NewState)
    end.

loop(State) ->
    receive
        {apply, Mod, Fun, Args} ->
            NewState = apply(Mod, Fun,
                             Args),
            loop(NewState)
    end.
```

- Matching only on unbound variables can be problematic
 - The pattern **{Mod, Fun, Arg}** can also match **{system, From, Req}** and break the server that expects other values
- Tag the messages to be sent and received
 - By tagging all messages and maintaining a functional interface, the matching is generally safer



Concurrency: using references

```
call(Pid, Message) ->
    Ref = make_ref(),
    Pid ! {request, {Ref, self()},
          Message},
    receive
        {reply, Ref, Reply} ->
            Reply
    end.

reply({Ref, Pid}, Reply) ->
    Pid ! {reply, Ref, Reply}.
```

- Similar requests and responses might come from different processes
- Multiple communications with replies can become confusing
- References help uniquely tag a set of messages to identify them



Concurrency: using references

```
call(Pid, Message) ->
  Ref = erlang:monitor(process, Pid),
  Pid ! {request, {Ref, self()}, Message},
  receive
    {reply, Ref, Reply} ->
      erlang:demonitor(Ref, [flush]), Reply
    {'DOWN', Ref, process, Pid, _} -> error(noproc)
  end.

reply({Ref, Pid}, Reply) -> Pid ! {reply, Ref, Reply}.
```

- Monitors allow the same properties than using **make_ref/0**
- They allow to find failures or missing processes earlier
- They must be taken down after a successful call



Conventions: nesting

- Avoid deeply nested code
- Keep only two levels of indentation for:
 - case
 - if
 - receive
 - funs
- Reduce indentation by
 - Pattern matching
 - creating temporary composite data types
 - Adding more functions



Conventions: nesting example

```
Case Oper of
  enabled ->
    case Admin of ->
      enabled -> restart(Board);
      disabled -> ok
    end;
  disabled ->
    ok
end.

case {Oper, Admin} of
  {enabled, enabled} -> restart(Board);
  {_, _} -> ok
end.
```



Conventions: code size

- Avoid long lines
 - Line should not be longer than 80 characters by convention
 - People who review the code by printing it have it easier
 - Can read more files with a split screen
- Avoid large modules
 - A module should usually not contain more than 500 lines of code, plus comments.
 - Not a strict metric, can be broken, but usually helps restricting each module to one responsibility



Conventions: code size

- Make your functions short
 - A function should rarely be above 50 lines and it should not span many pages
 - Versions of Erlang prior to R15 do not give line numbers in stack traces, makes it easier to debug
 - debugging trace flags can be enabled on specific functions only



Conventions: strategies

- Minimise the number of functions with side effects
- Abstract out common design patterns
- Avoid defensive programming
 - When you do not know how to handle a specific error, **let it crash**.
- Avoid copy-and-paste programming
- Don't comment out dead code, just delete it
 - Get it back from your repository if you need it again in the future



Style and Efficiency

- Style
- Efficiency
 - Pattern Matching
 - Recursion
 - IO Lists, Strings and Binaries
 - Garbage Collection
 - Concurrency



Pattern Matching: rearranging clauses

```
day(monday)    -> weekday;  
day(tuesday)   -> weekday;  
day(wednesday) -> weekday;  
day(thursday)  -> weekday;  
day(friday)    -> weekday;  
day(Name) when Name==saturday; Name==sunday -> weekend.
```

- There is nothing to gain by rearranging most of the clauses
- if the function is called more often with **friday**, it is not worth it to put **day(friday)** first before all other clauses.
- The compiler knows how to read clauses in the most efficient way possible without losing the semantics of the code:
 - A single instruction tries the first five clauses with a binary search
 - The last clause is tried if none of the first did match



Pattern Matching: exceptions

```
day(monday)    -> weekday;  
day(tuesday)   -> weekday;  
day(wednesday) -> weekday;  
day(Name) when Name==saturday; Name==sunday -> weekend;  
day(thursday)  -> weekday;  
day(friday)    -> weekday.
```

- The compiler does a fast search to know if the value matches **monday**, **tuesday**, or **wednesday**.
- If they do not match, the clause with **Name** then always matches
 - The guard test is executed
- If the guard test fails, the two last clauses are searched



Pattern Matching: exceptions

- Functions pattern matching with binaries are never rearranged
 - usually, placing the clause with the empty binary (**<<>>**) last is faster
 - This is version-specific and could change with newer releases
- Complex patterns and guards can make things harder for the compiler and lead for more manual adjustments and rearranging



Recursion

- Tail-recursive functions are not necessarily faster than non-tail-recursive functions
 - If the function generates a single value (ex.: factorial), tail-recursive functions are generally faster
 - If the function takes a list and returns a list of the same size, non-tail-recursive functions might be faster. ex.: **lists:map/2** is implemented in a non-tail-recursive manner
 - If the function is a server loop that never returns, you **must** use tail-recursive calls



IO Lists, Strings and Binaries

- Strings are implemented as lists of integers
 - They take more space
 - Useful if the string needs to be transformed and manipulated
- Binary strings are implemented as binaries
 - Compact with O(1) read access
 - Useful for storage, transport and limiting the size in memory



IO Lists, Strings and Binaries

```
1> IoList = [{"", "hello", " ", <<"hello",65>>, [<<$.,$ ,
$W>>, [{"}] | "rld"].
[34,"hello", " ",<<"helloA">>,<<". W">>,"o"],114,108,100]
2> io:format("~s~n",[IoList]).
"hello helloA. World
ok
```

- IO lists are lists of binaries, strings, or characters used for input/output operations
 - Automatically flattened by any **io** function, socket or port. No need to flatten it yourself.
 - They allow to append data in constant time by just wrapping the elements in a list: **[OldString | NewEndOfString]**.
 - Allows to mix the use of any string representation for a single element



IO Lists, Strings and Binaries

- For general string manipulation, use the **re** module
 - PCRE engine in C, faster than the obsolete **regexp** module
 - Works with all types of strings and can give the output format of your choice (index, list or binary)
- Use the **binary** module for binary string manipulation
- Use the **string** module for lists-based string manipulation



Garbage Collection

- Garbage Collection costs are largely proportional to the amount of live data on the heap
 - Large heaps for selected processes can give significant performance gains
- Try to avoid building large heap structures
 - **ets:foldl/3**, **ets:foldr/3**, **ets:tab2list/1**, etc.
 - Copying and rewriting lists
 - Use iolists rather than flattening strings when possible



Garbage Collection: heap size

- Set each process' heap size with
 - **spawn_opt(Module, Function, Args, OptionList)**
 - where **OptionList** includes **{min_heap_size, Size}**.
- Erlang heap memory is divided into the **old heap** and the **new heap**
 - Data in the new heap that survives a garbage collection sweep is moved to the old heap.
 - The option **{fullsweep_after, Number}** makes it possible to specify the number of garbage collections which should occur before the old heap is swept



Garbage Collection: heap size

- The **fullsweep_after** and **min_heap_size** options can be set globally for new processes by calling **erlang:system_flag(Flag, Value)**.
 - Use **fullsweep_after** only if you know there are problems with the memory consumption of your processes
 - the **min_heap_size** can be set when starting the VM with the **+h Size** flag
- You can force a garbage collection with **garbage_collect()** for the current calling process or **garbage_collect(Pid)** for a particular one



Concurrency: hibernation

```
erlang:hibernate(Module, Function, Arguments)
```

- If a particular process has to wait for a message for a while doing nothing, it can be compacted by calling **erlang:hibernate/3**. The process wakes up when receiving a message and keeps running with **Module:Function(Arguments)**.
 - Discards the call stack for the process
 - The process is garbage collected
 - All of the data is stored in one continuous heap, which is then shrunk to the exact same size as the data it holds.
- Hibernation is a tradeoff between processor use and memory use. Mostly useful for processes that are mostly inactive.



Concurrency: **message passing**

- Data is fully copied between processes
- Send messages about what happened, not the new state to carry
 - This tends to reduce message size, although this is generally a minor efficiency issue.
- Binaries larger than 64 bytes are passed around as pointers and are not copied between processes on a single node
 - Storing text or data that needs to be forwarded between many processes could be planned as binaries if required



Coding Strategies

First, make it work

Then, make it beautiful

**Then, if you really have to,
make it fast**



Style and Efficiency

- Style
- Efficiency
 - Pattern Matching
 - Recursion
 - IO Lists, Strings and Binaries
 - Garbage Collection
 - Concurrency

